References

[1] DC@KTH, *Feature Store for ML*, (accessed September 30, 2020), https://www.featurestore.org/.

[2] Logical Clocks, *Hopsworks Feature Store*, (accessed September 15, 2020), https://uploads-ssl.webflow.com/5e6f7cd3ee7f51d539a4da0b/5ef397dce942172c90259858_feature%20management%20platform%20v02.pdf.

[3] FEAST, *Architecture*, (accessed September 15, 2020), https://docs.feast.dev/user-guide/architecture.

[4] W. Pienaar, "Feast: Feature store for machine learning," 2019, anthill Inside 2019 https://hasgeek.com/anthillinside/2019/proposals/feast-feature-store-for-machine-learning-HBMmchibhze5ZpWqzhcian.

[5] J. Hermann, *Michelangelo - Machine Learning @Uber*, Mar 23, 2019, qCon https://www.infoq.com/presentations/uber-ml-michelangelo/?utm_source=youtube&utm_medium=link&utm_campaign=qcontalks.

[6] J. Hermann and M. D. Balso, *Meet Michelangelo: Uber's Machine Learning Platform*, 2017 (accessed September 30, 2020), https://eng.uber.com/michelangelo-machine-learning-platform/.

[7] V. Zanoyan and E. Shapiro, "Zipline—airbnb's declarative feature engineering framework," Okt 21, 2019, databricks https://www.youtube.com/watch?v=iUnO4MLAGDU&feature=emb_title&ab_channel=Databricks.

[8] N. Sarwar, "Operationalizing machine learning—managing provenance from raw data to predictions," 2018, databricks https://www.youtube.com/watch?v=iUnO4MLAGDU&feature=emb_title&ab_channel=Databricks.

[9] R. Romano, "Overview of wix's machine learning platform," Feb 18, 2020, wix Engineering Tech Talks https://www.youtube.com/watch?v=E8839ENL-WY&feature=emb_title&ab_channel=WixEngineeringTechTalks.

[10] N. Jain, "Real time machine learning inference platform," Nov 27, 2020, hagsgeekhttps://hasgeek.com/rootconf/2020-delhi/proposals/real-time-machine-learning-inference-platform-zoma-Jpk3vdwjcrQKGs19mGLibj.

[11] S. Khadder, "Feature stores: Building machine learning infrastructure on apache pulsar," Jun 24, 2020, pulsar Summit 2020 https://www.youtube.com/watch?v=eLJk5sOME0o&feature=emb_title&ab_channel=StreamNative.

[12] S. Canchi and T. Wenzel, *Managing ML Models @ Scale - Intuit's ML Platform*, Jul. 2020, https://www.usenix.org/conference/opml20/presentation/wenzel.

[13] S. Bathini, "Sony interactive entertainment: Powering playstation personalization to millions," Jun 17, 2020, aerospike SUMMIT '20 https://www.aerospike.com/resources/videos/summit20/sony-powering-playstation-personalization-to-millions/.

[14] J. Han, E. Haihong, G. Le, and J. Du, "Survey on nosql database," in *2011 6th international conference on pervasive computing and applications*. IEEE, 2011, pp. 363–366.

[15] A. Makris, K. Tserpes, V. Andronikou, and D. Anagnostopoulos, "A classification of nosql data stores based on key design characteristics." in *Cloud Forward*, 2016, pp. 94–103.

[16] G. Belalem, H. Matallah, and K. Bouamrane, "Evaluation of nosql databases: Mongodb, cassandra, hbase, redis, couchbase, orientdb," *International journal of software science and computational intelligence*, vol. 12, no. 4, pp. 71–91, 2020.

[17] N. Jatana, S. Puri, M. Ahuja, I. Kathuria, and D. Gosain, "A survey and comparison of relational and non-relational database," *International Journal of Engineering Research & Technology*, vol. 1, no. 6, pp. 1–5, 2012.

[18] F. Gessert and N. Ritter, "Scalable data management: Nosql data stores in research and practice," in *2016 IEEE 32nd International Conference on Data Engineering (ICDE)*. IEEE, 2016, pp. 1420–1423.

[19] A. Fox, S. D. Gribble, Y. Chawathe, E. A. Brewer, and P. Gauthier, "Cluster-based scalable network services," in *Proceedings of the sixteenth ACM symposium on Operating systems principles*, 1997, pp. 78–91.

[20] M. Kleppmann, *Designing Data-Intensive Applications: The Big Ideas Behind Reliable, Scalable, and Maintainable Systems*. Sebastopol: O'Reilly Media, Incorporated, 2017. ISBN 1449373321

[21] T. Weiss, *MySQL acquiring data management system vendor Alzato*, Oct 2003, https://www.computerworld.com/article/2572959/mysql-acquiring-data-management-system-vendor-alzato.html.

[22] R. Cattell, "Scalable sql and nosql data stores," *Acm Sigmod Record*, vol. 39, no. 4, pp. 12–27, 2011.

[23] Oracle, *Chapter 4 MySQL NDB Cluster Connector for Java*, (accessed Feb 5, 2021), https://dev.mysql.com/doc/ndbapi/en/mccj.html.

[24] B. Ocklin, "Sql faster than nosql, mysql ndb 8.0," Jan 31, 2020, oracle https://www2.slideshare.net/ocklin/mysql-ndb-cluster-80-sql-faster-than-nosql.

[25] Aerospike, *Aerospike Technical Documentation*, 2020 (accessed October 2, 2020), https://www.aerospike.com/docs/#.

[26] Aerospike, *Database Product Matrix*, 2020 (accessed October 2, 2020), https://www.aerospike.com/products/product-matrix/.

[27] D. Gembalczyk, F. M. Schuhknecht, and J. Dittrich, "An experimental analysis of different key-value stores and relational databases," *Datenbanksysteme für Business, Technologie und Web (BTW 2017)*, 2017.

[28] K. Kingsbury, "Redis-raft 1b3fbf6," Jun 23, 2020 (accessed October 1, 2020), aphyr https://jepsen.io/analyses/redis-raft-1b3fbf6.

[29] ——, "Jepsen: Redis," 2013 (accessed October 1, 2020), aphyr https://aphyr.com/posts/283-call-me-maybe-redis.

[30] Redis, *Redis Cluster Specification*, (accessed September 30, 2020), https://redis.io/topics/cluster-spec.

[31] Google Cloud, *Overview of Memorystore for Redis*, (accessed October 2, 2020), https://cloud.google.com/memorystore/docs/redis/redis-overview.

[32] G. DeCandia, D. Hastorun, M. Jampani, G. Kakulapati, A. Lakshman, A. Pilchin, S. Sivasubramanian, P. Vosshall, and W. Vogels, "Dynamo," *Operating systems review*, vol. 41, no. 6, p. 205, 2007.

[33] riak, *Riak KV 2.2.3*, (accessed October 2, 2020), https://docs.riak.com/riak/kv/2.2.3/.

[34] F. Chang, J. Dean, S. Ghemawat, W. C. Hsieh, D. A. Wallach, M. Burrows, T. Chandra, A. Fikes, and R. E. Gruber, "Bigtable: A distributed storage system for structured data," *ACM Transactions on Computer Systems (TOCS)*, vol. 26, no. 2, pp. 1–26, 2008.

[35] Apache HBase, *Welcome to Apache HBase ™*, (accessed October 5, 2020), http://hbase.apache.org/.

[36] Netflix, *Introduction*, (accessed September 19, 2020), https://hollow.how/.

[37] R. Srinivas, *Q&A with Drew Koszewnik on a Disseminated Cache, Netflix Hollow*, Dec 2016, https://www.infoq.com/news/2016/12/announcing-netflix-hollow/.

[38] B. F. Cooper, A. Silberstein, E. Tam, R. Ramakrishnan, and R. Sears, "Benchmarking cloud serving systems with ycsb," in *Proceedings of the 1st ACM symposium on Cloud computing*, 2010, pp. 143–154.

[39] B. Cooper, *YCSB*, (accessed September 30, 2020), https://github.com/brianfrankcooper/YCSB.

[40] K. Kingsbury, *Jepsen*, (accessed September 30, 2020), https://github.com/jepsen-io/jepsen.

[41] ——, *Jepsen: Aerospike*, 2015 (accessed October 1, 2020), aphyr https://aphyr.com/posts/324-jepsen-aerospike.

[42] ——, "Jepsen: Aerospike," 2018 (accessed September 30, 2020), aphyr https://jepsen.io/analyses/aerospike-3-99-0-3.

[43] ——, "Jepsen 9: A fsyncing feeling," May 8, 2018 (accessed October 1, 2020), gOTO 2018 https://www.youtube.com/watch?v=tRc0O9VgzB0&ab_channel=GOTOConferences.

[44] The Apache Software Foundation, "Apache beam: An advanced unified programming model," (accessed January 19, 2021), https://beam.apache.org/.

[45] Spotify, “Scio,” 2020 (accessed January 19, 2021), https://spotify.github.io/scio/.